

Operating Systems Notes

Marie Burer

16-1-24

Contents

1	Introduction to OS	5
2	OS Structure and Architecture	8
3	Process API and Abstraction	13
4	Proper Mechanism of Execution	18
5	Scheduling Policies	22
6	Introduction to Virtual Memory	24
7	Demand Paging	26
8	LRU Approximation	27
9	Stuff I didn't write down yet	27
10	Locks	28
11	Condition Variables	29
12	Machine Learning	29

Course Outline

Information:

- CS:3620 Operating Systems
- Instructor: Guanpeng Li
- Location: 110 MacLean Hall
- Hours: Tuesday/Thursday, 8:00 AM - 9:15 AM
- Prerequisite: CS:2210, CS:2230, and CS:2630 with a minimum grade of C-
- TA: A K M Muhitul Islam

Modality:

- Lectures in person
- Office Hours: Tuesday/Thursday, 2:30 PM - 4:30 PM
- Location: IATL 340

Grading:

- Homework: 40%
 - 6-8 assignments
 - Programming tasks

- In-class activities: 20%
 - Quizzes:
 - * Multiple-choice
 - * Close book
 - Exercises:
 - * Open questions
 - * Individual or group
 - * BYO laptop
- Final Exam: 40%
 - Technical interview
 - Written exam

Policy:

- 3 days late policy
- Cannot make up quizzes or exercises
- Can discuss homework with colleagues
- Cannot share code
- Cannot copy from internet
- If you copy code from StackOverflow, credit it
- When in doubt, ask the instructor/TA

1 Introduction to OS

What is an operating system?

An operating system (OS) is system software that manages computer hardware and software resources and provides common services for computer programs. (Wikipedia)

- Definition has changed over years
 - Originally, very bare bones
 - Evolved to include more
- Operating System (OS)
 - Interface between the user and the architecture
 - Implements a virtual machine that is easier to program than raw hardware
- Middleware between user programs and system hardware
- Manages hardware: CPU, main memory, I/O devices

OS: Traditional View:

- Interface between user and architecture
 - Hides architectural details
- Implements Virtual machine:

- Easier to program
- Illusionist
 - Bigger, faster, reliable
- Government
 - Divides resources
 - "Taxes" = overhead

New Developments in OS:

- Operating systems: active field of research
 - Demands on OS's growing
 - New application spaces (Web, Grid, Cloud)
 - Rapidly evolving hardware
- Open-source operating systems
 - Linux etc.
 - You can contribute to and develop OS's
 - Excellent research platform

OS: Salient Features

- Services: The OS provides standard services which the hardware implements

- Coordination: The OS coordinates multiple applications and users to achieve fairness and efficiency
- Design an OS so that the machine is convenient to use and efficient

Why study OS?

- Abstraction: How to get the OS to give users an illusion of infinite resources
- System Design: How to make tradeoffs between performance, convenience, simplicity, and functionality
- Basic Understanding: The OS provides the services that allow application programs to work at all
- System Intersection Point: The OS is the point where hardware and software meet

Not many operating systems are under development, so you are unlikely to get a job building an OS. However, understanding operating systems will enable you to use your computer more efficiently

Build Large Computer Systems

- OS as an example of large system design
- Goals:

2 OS Structure and Architecture

Modern OS Functionality:

- Concurrency
- I/O Devices
- Memory management
- Files
- Distributed systems and networks

OS Principles:

- OS as juggler
- OS as government
- OS as complex system
- OS as history teacher

Protection:

- Kernel mode vs User Mode

To protect the system from aberrant users, some instructions are restricted to use only by the OS

Users may not

- address I/O directly
- use instructions that manipulate the state of memory
- set the mode bits that determine user or kernel mode
- disable and enable interrupts
- halt the machine

In kernel mode, the OS can do all of these things

- Hardware needs to support at least both kernel and user

Crossing Protection Boundaries:

System call: OS procedure that executes privileged instructions

- Causes a trap
- The trap handler uses the parameter to jump
- The handler saves caller's state so it can restore control
- The architecture must permit this

Memory Protection:

- Architecture must provide support to the OS to protect user programs and itself
- The simplest technique is to use base and limit registers
- Instantiated by the OS before starting a program

- The CPU checks user reference ensuring it falls between base and limit values

Memory Hierarchy:

Higher = small, fast, more \$, less latency

- registers (1 cycle)
- L1 (2 cycle)
- L2 (7 cycle)
- RAM (100 cycle)
- Disk (40,000,000 cycle)
- Network (200,000,000+ cycle)

Process Layout in Memory

- Stack
- Gap
- Data
- Text

Caches

- Access to main memory is expensive

- Caches are small, fast, inexpensive
 - Different sizes and locations:
 - * L1: on-chip, smallish
 - * L2: next to chip, larger
 - * L3: pretty large, on bus

Traps

- Traps: special conditions detected by architecture
- On detecting a trap, the hardware
 - Saves the state of the process
 - Transfers control to appropriate trap handler
 - * The CPU indexes the memory-mapped trap vector with the trap number,
 - * then jumps to the address given in the vector, and
 - * starts to execute at that address
 - * On completion, the OS resumes execution of the process

I/O Control

- Each I/O device has a little processor that enable autonomous function
- CPU issues commands to I/O devices
- When an I/O device complete, it issues an interrupt

- CPU stops and processes the interrupt

Memory Mapped I/O

- Enable direct access to I/O controller
- PCs reserve a part of the memory and put the device manager in that memory
- Access becomes very fast

Interrupt-based asynchronous I/O

- Device controller has its own small processor
- Puts an interrupt on the bus when done
- CPU gets interrupt
 - Save critical state
 - Disable interrupts
 - Save state that interrupt handler will modify
 - Invoke interrupt handler using the interrupt vector
 - Restore software state
 - Enable interrupts
 - Restore hardware state

3 Process API and Abstraction

OS provides process abstraction

- When you run an exe file, the OS creates a process = a running program
- OS timeshares the CPU across multiple processes: virtualizes CPU
 - No. of running processes \leq No. of physical CPU cores
 - Programs are responsive to users, even when the CPU is busy
- OS has a CPU scheduler that picks one of the active processes to execute
 - Policy: which to run
 - Mechanism: how to context switch

What constitutes a process?

- PID
- Memory image
 - Code and data (static)
 - Stack and heap (dynamic)
- CPU context: registers
 - Program counter

- Current operands
 - Stack pointer
- File descriptors
 - Pointers to open files and devices
 - e.g. STDOUT, STDIN, STDERR

States of a process

- Running: executing
- Ready: waiting
- Blocked: not ready to run
 - Why? Waiting for something
 - Unblocked when disk issues an interrupt
- New: being created

OS data structures

- OS maintains a data structure of all active processes
- Information about each process is stored in a process control block
 - Identifier
 - State

- Pointers to related processors
- CPU context
- Memory pointers
- Pointers to open files

What API does the OS provide to user programs?

- API
- Set of system calls
 - System call is a function call into the OS code that runs at a higher privilege
 - Sensitive operations are allowed only at this high privilege
 - Some blocking system calls cause the process to be blocked and descheduled

Should we rewrite programs for each OS?

- POSIX API: a standard set of system calls that an OS must implement
 - POSIX: Portable Operating System Interface
 - Most modern OSes are POSIX compliant
- Program language libraries hide the system calls

Process related system calls in Unix

- `fork()` creates a child process
- `exec()` makes a process execute a given executable
- `exit()` terminates a process
- `wait()` causes a parent to block until child terminates
- Many variants exist with different arguments

What happens during a fork?

- A new process is created by copying the parent's memory image
- The new process is added to process list and schedules
- Then the parent and child operate independently
- On success, the PID of the child process is returned in the parent
 - 0 is returned in the child
 - On failure, -1 is returned to the parent and no child process is created

What happens during `exec`?

- After fork, parent and child are running same code
Not too useful!
- A process can run `exec()` to load a new process

How does a shell work?

- In a basic OS, the init process is created after initialization of hardware
- The init process spawns a shell like bash
- Shell reads user command, forks a child, execs the command executable, waits, and reads
- Common commands like ls are all executables that are simply exec'd by the shell

4 Proper Mechanism of Execution

Low-level mechanisms

- How does the OS run a process?
- How does it handle a system call?
- How does it context switch?

Process Execution

- OS allocates memory and creates memory image
- Points CPU program counter to current instruction
- After setup, OS steps out of the way

A simple function call

- A function call translates to a jump instruction
- A new stack frame is pushed to stack and stack pointer is updated
- Old value of PC is saved to stack frame and PC is updated
- Stack frame contains return value, function arguments etc.

How is a system call different?

- System calls are in kernel memory space
- Kernel uses a separate stack from the user stack

- Kernel does not trust user provided addresses
- Instead, it sets up an Interrupt Descriptor Table to keep track

Mechanism of system call: trap instruction

- When a system call must be made, a special trap instruction is run
- Trap execution
 - Move CPU to higher privilege level
 - Switch to kernel stack
 - Save context on kernel stack
 - Look up address in IDT and jump to trap handler function in OS code

More on the trap instruction

- Executed on hardware in following cases:
 - System call (program needs OS service)
 - Program fault (illegal stuff)
 - Interrupt (device needs attention)
- Across all cases, the mechanism is the same
- IDT has many entries: which to use?

- System calls and interrupts store a number in a CPU register before calling trap, as an IDT index

Return from trap

- When OS is done handling, it calls a special return-from-trap instruction
 - Restore context
 - Lower privilege
 - Restore PC
- OS doesn't always return to the same process

Why switch?

- Sometimes when OS is in kernel mode, it cannot return back to the same process because of terminating or a blocking system call
- Sometimes it doesn't want to (timeshare)
- In such cases, a context switch is done

The OS scheduler

- Two parts
 - Scheduler which picks processes
 - Mechanism to switch to that process

- Non preemptive schedulers are polite (only if process blocked or terminated)
- Preemptive schedulers can switch even when a process is ready to continue
 - CPU generates periodic timer interrupts
 - After servicing, the OS checks if the process has been running too long

5 Scheduling Policies

What are we trying to optimize?

- Maximize utilization
- Minimize average turnaround time
- Minimize average response time
- Fairness
- Minimize overhead

FIFO

- Example: three processes arrive at $t=0$ in the order A,B,C
- Schedule A, then B, then C
- Problem: convoy effect (if the first process is long, turnaround times go high)

SJF (Shortest Job First)

- Provably optimal when all processes arrive together
- SJF is non-preemptive, so short jobs can still get stuck when short jobs arrive late

STCF (Shortest Time-to-Completion)

- Pre-emptive scheduler
- Pre-empt running task if time left is more than that of new arrival

RR (Round Robin)

- Every process executes for a fixed quantum slice
- Slice big enough to amortize cost of context switch
- Pre-emptive
- Good for response time and fairness
- Bad for turnaround time

Schedulers in real systems

- Real schedules are more complex
- For example, Linux uses a Multi Level Feedback Queue (MLFQ)
 - Many queues, with a priority order
 - Process from highest priority queue scheduled first
 - Within same priority, any algorithm like RR
 - Priority of process decays with its age

6 Introduction to Virtual Memory

Why virtualize?

- It's messy!
- Earlier, memory only had code of one running processes
- Now, multiple active processes timeshare
- Need to hide this complexity from user

Abstraction: Virtual Address space

- Virtual address space: every process assumes it has maximum memory space
- contains program code, heap, and stack
- CPU issues loads and stores to virtual addresses

How is actual memory reached?

- Address translation from VA to PA
- OS allocates memory and tracks location of processes
- Actual translation by MMU

Example: Paging

- OS divides virtual address space into fixed size pages, physical memory into frames

- To allocate memory, a page is mapped to a free physical frame
- Page table stores mappings from virtual page to physical frame
- MMU has access to page tables

Goals of memory virtualization

- Transparency: user programs shouldn't handle the details
- Efficiency: minimize overhead
- Isolation and protection: a user program should not be able to access anything outside its address space

How can a user allocate memory?

- OS allocates a set of pages to the memory image of the process
- Within this image, the process can allocate memory

7 Demand Paging

Main memory is not always enough

We have x amount of frames in physical memory, but n amount of frames in swap space.

Swap these on demand

Present bit in PTE, indicates if a page is in memory or not.

If present bit is not set on MMU lookup, page fault

OS moves to kernel mode, switches to another process while lookup happens

What happens on memory access:

- CPU issues load to VA for data
- Checks all 3 caches
- In case of miss, goes to main memory
- MMU issues request to TLB
- If no hit, page fault

Complications during a page fault:

- If no room in memory, must issue a swap.
- Time consuming

Replacement policies: FIFO, LRU

8 LRU Approximation

Second chance algorithm:

Maintain a reference bit next to each frame in a circular linked list

If the reference bit is 1, set it to 0 and move on

If it is 0, replace

If we use the page, set the bit to 1

Enhanced second chance, two bits

9 Stuff I didn't write down yet

10 Locks

Locks: Basic idea

When updating some shared variable, we can use a lock variable to protect it

- Mutual exclusion
- Fairness
- Low overhead

Naive implementation:

Lock disables interrupts

Unlock enables them

Flaws:

- Disable and enable is privileged and can be misused
- Doesn't work for multi-processor

Next implementation:

Lock spins until a flag is 0, then updates it to 1 after

Unlock sets to 0

Flaws:

- Busily waiting, takes CPU to do nothing
- Still has race conditions (FAILED)

Software implementation: almost impossible

Modern architecture uses hardware atomic instructions

Atomic instruction: `testAndSet`

Delivers mutual exclusion and fairness

Alternative to spinning: `yield`, gives up CPU

11 Condition Variables

Another type of synchronization

Locks allow mutual exclusion

Another common requirement is waiting and signaling

Inefficient to busy-wait, we need a new synchronization primitive.

12 Machine Learning

Stuff I didn't write down

We can also have multiple neurons, and multiple layers of neurons